



Shared Memory – OpenMP

Computer Architecture and Operating Systems Department

Universidad Aut3noma de Barcelona

tomas.margalef@uab.es



Shared Memory

- OpenMP is currently the programming standard for the shared memory model on multicore systems.



Shared Memory

- Thread based model.
- Threads read and write shared variables.
- Synchronization mechanism are offered.
- It is possible to change the attributes of threads and data for minimizing synchronization.



OpenMP



- OpenMP is not an automatic parallelization tool:
 - Programmers must specify parallelism explicitly.
- OpenMP is not only for exploiting loop parallelism:
 - It also offers functionalities for other forms of parallelism.



OpenMP

- OpenMP is not a programming language:
 - It is structured as extensions using directives to base languages like Fortran or C.
- OpenMP is not only a research project:
 - Many commercial compilers support OpenMP.




OpenMP


- OpenMP is an API (Application Program Interface) for parallel programming in shared memory systems.
- It's main objective is to easily parallelize existing applications.
- It's based in directives introduced in the program as special comments (pragmas).

OpenMP Execution Model

- Initially based on a FORK-JOIN model



Shared Memory Model



- Data can be Shared or private.
- All threads have access to the same globally Shared memory.
- Shared data is accessible by all threads.
- Private data can be accessed only by the threads that owns it.
- Data transfer is transparent to the programmer.
- Synchronization takes place, but it is mostly implicit.

Shared Memory Model

- In a Shared memory parallel program variables have a “*label*” attached to them:
 - Labelled “*Private*” ➡ Visible to one thread only.
 - Change made in local data, is not seen by others.
 - Labelled “*Shared*” ➡ Visible to all threads.
 - Change made in global data, is seen by all others.

Components of OpenMP

➤ Directives

- Parallel regions
- Work sharing
- Synchronization
- Data scope attributes
 - private
 - firstprivate
 - lastprivate
 - shared

➤ Environment variables

- Number of threads
- Scheduling type
- Dynamic thread adjustment
- Nested parallelism

➤ Runtime Environment

- Number of threads
- Dynamic thread adjustment
- Nested parallelism
- Timers
- API for locking



User Interface

- Compiler directives
 - There are control structures and data attributes structures.
 - Compilers ignore these directives (they are just comments) unless the proper options are used when compiling (“-mp” or “-fopenmp”).

User Interface

- Environment variables

- Set of variables for controlling some parameters

```
setenv OMP_NUM_THREADS 8
```

- Runtime Library

- Set of functions for controlling some parameters, such as the number of threads to be

```
omp_set_num_threads (128)
```

The omp parallel directive

- A parallel region is a block of code executed by multiple threads simultaneously.

```
#pragma omp parallel [clause[[,] clause] ...]  
{  
    "This will be executed in parallel by each thread"  
}  
    (implied barrier)
```

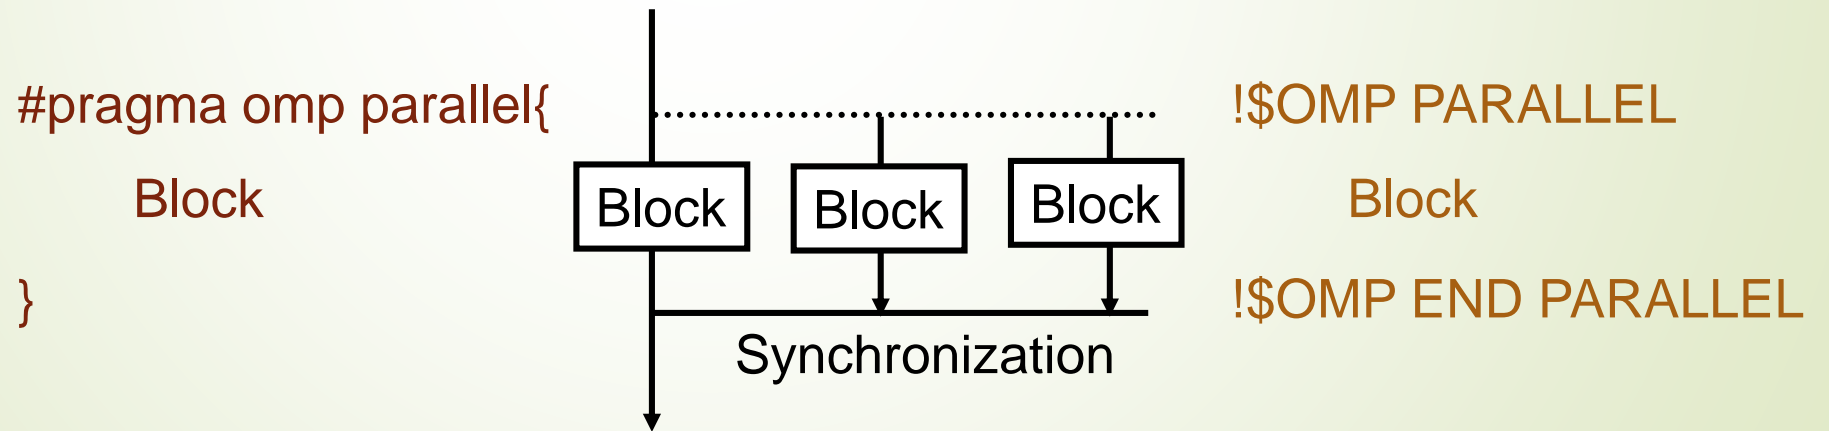
C

```
!$omp parallel [clause[[,] clause] ...]  
  
    "This will be executed in parallel by each thread"  
  
!$omp end parallel    (implied barrier)
```

Fortran

The omp parallel directive

- Define a parallel region.
- It does the “fork” and “join”.
- The number of threads is constant in the parallel region.



The omp parallel clauses

➤ omp parallel supports the following clauses:

- | | |
|----------------|---------------------|
| ➤ if | (scalar expression) |
| ➤ private | (list) |
| ➤ shared | (list) |
| ➤ default | (none/shared) |
| ➤ reduction | (operator: list) |
| ➤ firstprivate | (list) |
| ➤ num_threads | (scalar_int_expr) |

The omp for/do directive

- The iterations of the loop are distributed over the threads.

```
#pragma omp for [clause[,] clause] ...]
```

<original for-loop>

C

```
!$omp parallel [clause[,] clause] ...]
```

<original do-loop>

Fortran

- Clauses supported:
 - private firstprivate
 - reduction
 - schedule
 - nowait

The omp for/do directive

- `#pragma omp for / !$omp do - !$omp end do`
- It is for classical parallel loops.
- It must be in a parallel region.
- Loop iterations are distributed among available threads.
- Loop index is by default private to each thread.

The omp for directive - example

```
{  
  
    for (i=0; i<n-1; i++)  
        b[i] = (a[i] + a[i+1])/2;  
  
    for (i=0; i<n-1; i++)  
        d[i] = 1.0/c[i];  
  
}
```

The omp for directive - example

```
#pragma omp parallel default (none) shared (n,a,b,c,d) private (i)
{

    for (i=0; i<n-1; i++)
        b[i] = (a[i] + a[i+1])/2;

    for (i=0; i<n-1; i++)
        d[i] = 1.0/c[i];

}
```

The omp for directive - example


```
#pragma omp parallel default (none) shared (n,a,b,c,d) private (i)
{

    #pragma omp for nowait
    for (i=0; i<n-1; i++)
        b[i] = (a[i] + a[i+1])/2;

    #pragma omp for nowait
    for (i=0; i<n-1; i++)
        d[i] = 1.0/c[i];


}
```

Example: Matrix-Vector multiplication



```
for (i=0; i<m; i++)  
{  
    sum = 0.0;  
    for (j=0; j<n; j++)  
        sum += b[i][j]*c[j]  
    a[i] = sum;  
}
```

Example: Matrix-Vector multiplication



```
#pragma omp parallel for default (none) shared (m,n,a,b,c) private (i,j,sum)
for (i=0; i<m; i++)
{
    sum = 0.0;
    for (j=0; j<n; j++)
        sum += b[i][j]*c[j]
    a[i] = sum;
}
```


Example: Matrix-Vector multiplication

```
#pragma omp parallel for default (none) shared (m,n,a,b,c) private (i,j,sum)
for (i=0; i<m; i++)
{
    sum = 0.0;
    for (j=0; j<n; j++)
        sum += b[i][j]*c[j]
    a[i] = sum;
}
```

TID = 0

```
for (i = 0, 1, 2, 3, 4)
    i = 0
    sum =  $\sum b[0][j]*c[j]$ 
    a[0] = sum
    i = 1
    sum =  $\sum b[1][j]*c[j]$ 
    a[1] = sum
```

...

TID = 1

```
for (i = 5, 6, 7, 8, 9)
    i = 5
    sum =  $\sum b[5][j]*c[j]$ 
    a[5] = sum
    i = 6
    sum =  $\sum b[6][j]*c[j]$ 
    a[6] = sum
```

...

TID = 2

```
for (i = 10, 11, 12, 13, 14)
    i = 10
    sum =  $\sum b[10][j]*c[j]$ 
    a[10] = sum
    i = 11
    sum =  $\sum b[11][j]*c[j]$ 
    a[11] = sum
```

...


TID = 3

```
for (i = 15, 16, 17, 18, 19)
    i = 15
    sum =  $\sum b[15][j]*c[j]$ 
    a[15] = sum
    i = 16
    sum =  $\sum b[16][j]*c[j]$ 
    a[16] = sum
```

...


Example: Matrix-Vector multiplication

```
int m, n, i, j, sum  
int a[100][100], b[100], c[100]
```



Example: Matrix-Vector multiplication

```
#pragma omp parallel for default (none) shared (m,n,a,b,c) private (i,j,sum)
```



TID = 0

i

j

sum

TID = 1

i

j

sum

TID = 2

i

j

sum

TID = 3

i

j

sum

The default clause

- default (none | shared | private)
 - none
 - No implicit defaults.
 - Have to scope all variables explicitly.
 - shared
 - All variables are shared.
 - The default in absence of an explicit “default” clause.
 - private
 - All variables are shared.

The reduction clause

- reduction (operator : list)
 - Reduction variable(s) must be shared variables.
 - Operators:
 - +, -, *, max, min, ...

```
int maximum = 0, i;  
int Vector[1000];  
#pragma omp parallel for shared (Vector) private (i) reduction (max : maximum)  
{  
    for (i=0; i<1000; i++)  
        if (Vector[i] > maximum)  
            maximum = Vector[i];  
}
```

The ordered clause

- The iterations of the loop are executed in the same order than the sequential execution (1 thread at a time).

```
#pragma omp parallel for ordered
{
    for (i=0; i<n; i++)
        printf ("n = %d\\d, n);
}
```

Load balancing

- Load balancing is an important aspect of performance.
- For regular operations (e.g. a vector addition), load balancing is not an issue.
- For less regular workloads, care needs to be taken in distributing the work over the threads.
- Irregular workloads include transposing a matrix, multiplication of irregular matrices, ...
- The **schedule** clause supports various iteration scheduling algorithms.

The schedule clause

- `schedule (static | dynamic | guided [, chunk])` (runtime)
 - `static [, chunk]`
 - Distribute iterations in blocks of size “chunk” over the threads in a round-robin fashion.
 - In absence of “chunk”, each thread executes approx. N/P iterations.
 - `dynamic [, chunk]`
 - Fixed portions of work; size is controlled by the value of chunk.
 - The default in absence of an explicit “default” clause.

The schedule clause

- `schedule (static | dynamic | guided [, chunk]) (runtime)`
 - `guided [, chunk]`
 - Same dynamic behavior as “dynamic”, but size of the portion of work decreases exponentially.
 - `runtime`
 - Iteration scheduling is set at runtime through environment variable `OMP_SCHEDULE`.

The schedule clause

- `schedule (static | dynamic | guided [, chunk]) (runtime)`
 - `guided [, chunk]`
 - Same dynamic behavior as “dynamic”, but size of the portion of work decreases exponentially.
 - `runtime`
 - Iteration scheduling is set at runtime through environment variable `OMP_SCHEDULE`.

The schedule clause example

- 80 iterations on 4 threads

TID	0	1	2	3
Static No chunk	0-19	20-39	40-59	60-79
Static Chunk 10	0-9 40-49	10-19 50-59	20-29 60-69	30-39 70-79
Dynamic Chunk 10	0-9	10-19 40-49 60-69	20-29 70-79	30-39 50-59
Guided Chunk 10	0-19	20-34 57-66	35-46 67-76	47-56 77-79

The nowait clause


- To minimize synchronization, some OpenMP pragmas support the **nowait** clause.
- If present, threads will not synchronize/wait at the end of that particular construct.

```
#pragma omp parallel for default (none) shared (n,a,b) private (i) nowait  
{  
    for (i=0; i<n-1; i++)  
        b[i] = (a[i] + a[i+1])/2;  
}
```

The omp barrier directive

- All threads wait until the last arrives to the barrier

```
#pragma omp parallel num_threads (3)
{
    Block1;
    #pragma omp barrier
    Block2;
}
```



The omp sections directive

- Must be in a parallel region.
- Sections are distributed among threads.
- Each thread executes a different section.
- Allows task level parallelism.
- SECTION pragma defines each section.

The omp sections directive - example

```
#pragma omp parallel num_threads (3)
{
    #pragma omp sections
    {
        #pragma omp section {
            Block1;
        }
        #pragma omp section {
            Block2;
        }
        #pragma omp section {
            Block3;
        }
    }
}
```


The omp sections directive - example

```
#pragma omp parallel default (none) shared (n,a,b,c,d) private (i)
{
    #pragma omp sections nowait
    {
        #pragma omp section
        for (i=0; i<n-1; i++)
            b[i] = (a[i] + a[i+1])/2;
        #pragma omp section
        for (i=0; i<n-1; i++)
            d[i] = 1.0/c[i];
    }
}
```

The omp single directive

- The code included in the single section will be executed only by one thread.


```
#pragma omp parallel num_threads (3)
{
    Block1;
    #pragma omp single {
        Block2;
    }
    Block3;
}
```



The omp master directive

- The code included in the master section will be executed only by master thread.

```
#pragma omp parallel num_threads (3)
{
    Block1;
    #pragma omp master {
        Block2;
    }
    Block3;
}
```



The omp critical directive

- A critical section is a block of code which can be executed by only one thread at a time.
- Can be used to protect updates to shared variables.

```
sum = 0.0;
#pragma omp parallel for shared(sum, A, B) private (i)
{
    for (i=0; i<N; i++){
        A[i] = B[i]/2;
        #pragma omp critical
            sum += A[i];
    }
}
```

OpenMP 3.0

➤ TASK Construct

- The TASK construct defines an explicit task, which may be executed by the encountering thread, or deferred for execution by any other thread in the team.
- The data environment of the task is determined by the data sharing attribute clauses.

#pragma omp task *[clause ...]*

- The TASKWAIT construct specifies a wait on the completion of child tasks generated since the beginning of the current task.

#pragma omp taskwait

The task directive - example

```
struct node {node *next, .../*data*/}  
void process_list_items (node* head)  
{  
    #pragma omp parallel  
    {  
        #pragma omp single  
        {  
            for (node* p = head; p; p = p->next)  
            {  
                #pragma omp task firstprivate(p)  
                process(p);  
            }  
        }  
    }  
}
```

OpenMP 3.0

■ TASKGROUP Construct

- The TASKGROUP construct specifies a wait on completion of child tasks of the current task and their descendent tasks
- A TASKGROUP region binds to the current task region. The binding thread set of the taskgroup region is the current team

■ TASKYIELD Construct

- The TASKYIELD construct specifies that the current task can be suspended in favor of execution of a different task



OpenMP 4.0

- Support thread affinity policies (proc_bind, get_proc_bin, OMP_PLACES)
- Support execution on devices (accelerators)
(omp_set_default_device, omp_get_default_device, omp_get_num_devices, omp_get_num_teams, omp_get_team_num, and omp_is_initial_device)
- Reduction clause extended to support user defined reductions
- The concept of cancellation is added

Run-time Library

- `OMP_SET_NUM_THREADS (SCALAR)`
 - Sets the number of threads that will be used in the next parallel region.
 - Only works if called from a sequential portion of the program.
- `OMP_GET_NUM_THREADS ()`
 - Returns the number of threads in the parallel region where it's called.
 - The default number of threads depends on the application.
- `OMP_GET_THREAD_NUM ()`
 - Returns the thread id of the thread that calls it.
 - Master thread has id 0.

Run-time Library

- `omp_in_parallel`
- `omp_set_dynamic`
- `omp_get_dynamic`
- `omp_set_nested`
- `omp_get_nested`
- `omp_set_schedule`
- `omp_get_schedule`
- `omp_get_thread_limit`
- `omp_set_max_active_levels`
- `omp_get_max_active_levels`
- ...

References

- Ruud van der Pas “An Introduction Into OpenMP”. Sun Microsystems. iWOMP 2005.
- OpenMP Architecture Review Board. OpenMP Application Program Interface. Version 2.5. May 2005.
- OpenMP Application Program Interface. Version 4.0. July 2013
- www.openmp.org
- computing.llnl.gov/tutorials/openMP/